

ASSIGNMENT-2

Assignment Date	29 September 2022
Name	Parkavi.P.T
Rollnumber	820319205026
Maximum Marks	2Marks

```
import pandas as pd
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
sns.set_style('darkgrid')
```

```
sns.set(font_scale=1.3)
```

```
In [25]:
```

```
df=pd.read_csv("/content/drive/MyDrive/IBM/Assignment - 2 /Churn_Modelling.csv")
```

```
In [26]:
```

```
df.head ()
```

```
Out [26]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	Num
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3
3	4	15701354	Boni	699	France	Female	39	1	0.00	2

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	Num
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1

```
In [29]:
df.drop(["RowNumber","CustomerId","Surname"], axis=1, inplace=True)
```

```
In [30]:
df.info ()
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 10000 entries, 0 to 9999

Data columns (total 11 columns):

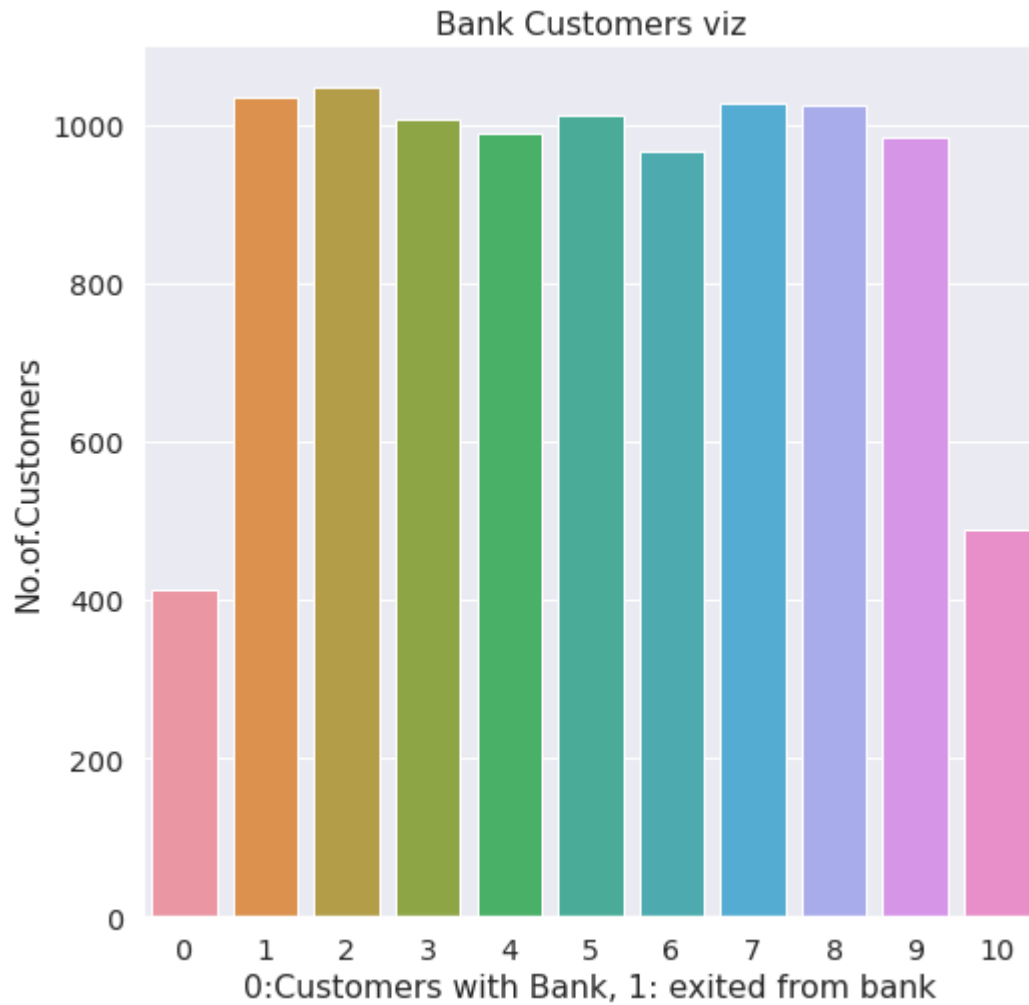
#	Column	Non-Null Count	Dtype
---	-----	-----	----
0	CreditScore	10000 non-null	int64
1	Geography	10000 non-null	object
2	Gender	10000 non-null	object
3	Age	10000 non-null	int64
4	Tenure	10000 non-null	int64
5	Balance	10000 non-null	float64
6	NumOfProducts	10000 non-null	int64
7	HasCrCard	10000 non-null	int64
8	IsActiveMember	10000 non-null	int64
9	EstimatedSalary	10000 non-null	float64
10	Exited	10000 non-null	int64

dtypes: float64(2), int64(7), object(2)

memory usage: 859.5+ KB

```
In [28]:
#Perform Univariate Analysis
plt.figure(figsize=(8,8))
sns.countplot(x='Tenure',data=df)
```

```
plt.xlabel('0:Customers with Bank, 1: exited from bank')
plt.ylabel('No.of.Customers')
plt.title("Bank Customers viz")
plt.show()
```




In[9]#Perform Univariate Analysis

```
plt.figure(figsize=(8,8))
sns.kdeplot(x=df['Balance'])
```

Out[9]:

<matplotlib.axes._subplots.AxesSubplot at 0x7fa0c03906d0>




```
In[10]#Perform Bivariate Analysis
```

```
plt.scatter(df.Age,df.Balance)
```

```
Out[10]:
```

```
<matplotlib.collections.PathCollection at 0x7fa0d35a7dd0>
```



#Perform Bivariate Analysis

df.corr()

Out[54]:

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsAct
CreditScore	1.000000	0.007888	-0.003965	0.000842	0.006268	0.012238	-0.005458	0.025
Gender	0.007888	1.000000	0.022812	0.003739	0.069408	0.003972	-0.008523	0.006
Age	-0.003965	0.022812	1.000000	-0.009997	0.028308	-0.030680	-0.011721	0.085
Tenure	0.000842	0.003739	-0.009997	1.000000	-0.012254	0.013444	0.022583	-0.02
Balance	0.006268	0.069408	0.028308	-0.012254	1.000000	-0.304180	-0.014858	-0.01
NumOfProducts	0.012238	0.003972	-0.030680	0.013444	-0.304180	1.000000	0.003183	0.009
HasCrCard	-0.005458	-0.008523	-0.011721	0.022583	-0.014858	0.003183	1.000000	-0.01

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsAct
IsActiveMember	0.025651	0.006724	0.085472	-0.028362	-0.010084	0.009612	-0.011866	1.000
EstimatedSalary	-0.001384	-0.001369	-0.007201	0.007784	0.012797	0.014204	-0.009933	-0.010
Exited	-0.027094	0.035943	0.285323	-0.014001	0.118533	-0.047820	-0.007138	-0.150

In [36]:

```

#Perform Bivariate Analysis

import statsmodels.api as sm

#define response variable
y = df['CreditScore']

#define explanatory variable
x = df[['EstimatedSalary']]

#add constant to predictor variables
x = sm.add_constant(x)

#fit linear regression model
model = sm.OLS(y, x).fit()

#view model summary
print(model.summary())

```

OLS Regression Results

=====			
Dep. Variable:	CreditScore	R-squared:	0.000

Model: OLS Adj. R-squared: -0.000
Method: Least Squares F-statistic: 0.01916
Date: Sat, 24 Sep 2022 Prob (F-statistic): 0.890
Time: 05:06:19 Log-Likelihood: -59900.
No. Observations: 10000 AIC: 1.198e+05
Df Residuals: 9998 BIC: 1.198e+05
Df Model: 1
Covariance Type: nonrobust

```
=====
              coef  std err          t    P>|t|    [0.025    0.975]
-----
```

```
const          650.7617    1.940    335.407    0.000    646.958    654.565
EstimatedSalary -2.326e-06  1.68e-05   -0.138    0.890   -3.53e-05    3.06e-05
```

```
=====
Omnibus:          132.939 Durbin-Watson:          2.014
Prob(Omnibus):          0.000 Jarque-Bera (JB):          84.242
Skew:           -0.072 Prob(JB):          5.10e-19
Kurtosis:          2.574 Cond. No.          2.32e+05

=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.32e+05. This might indicate that there are strong multicollinearity or other numerical problems.

/usr/local/lib/python3.7/dist-packages/statsmodels/tsa/tsatools.py:142: FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only

```
x = pd.concat(x[:,::order], 1)
```

In [35]:

```
#Perform Multivariate Analysis
```

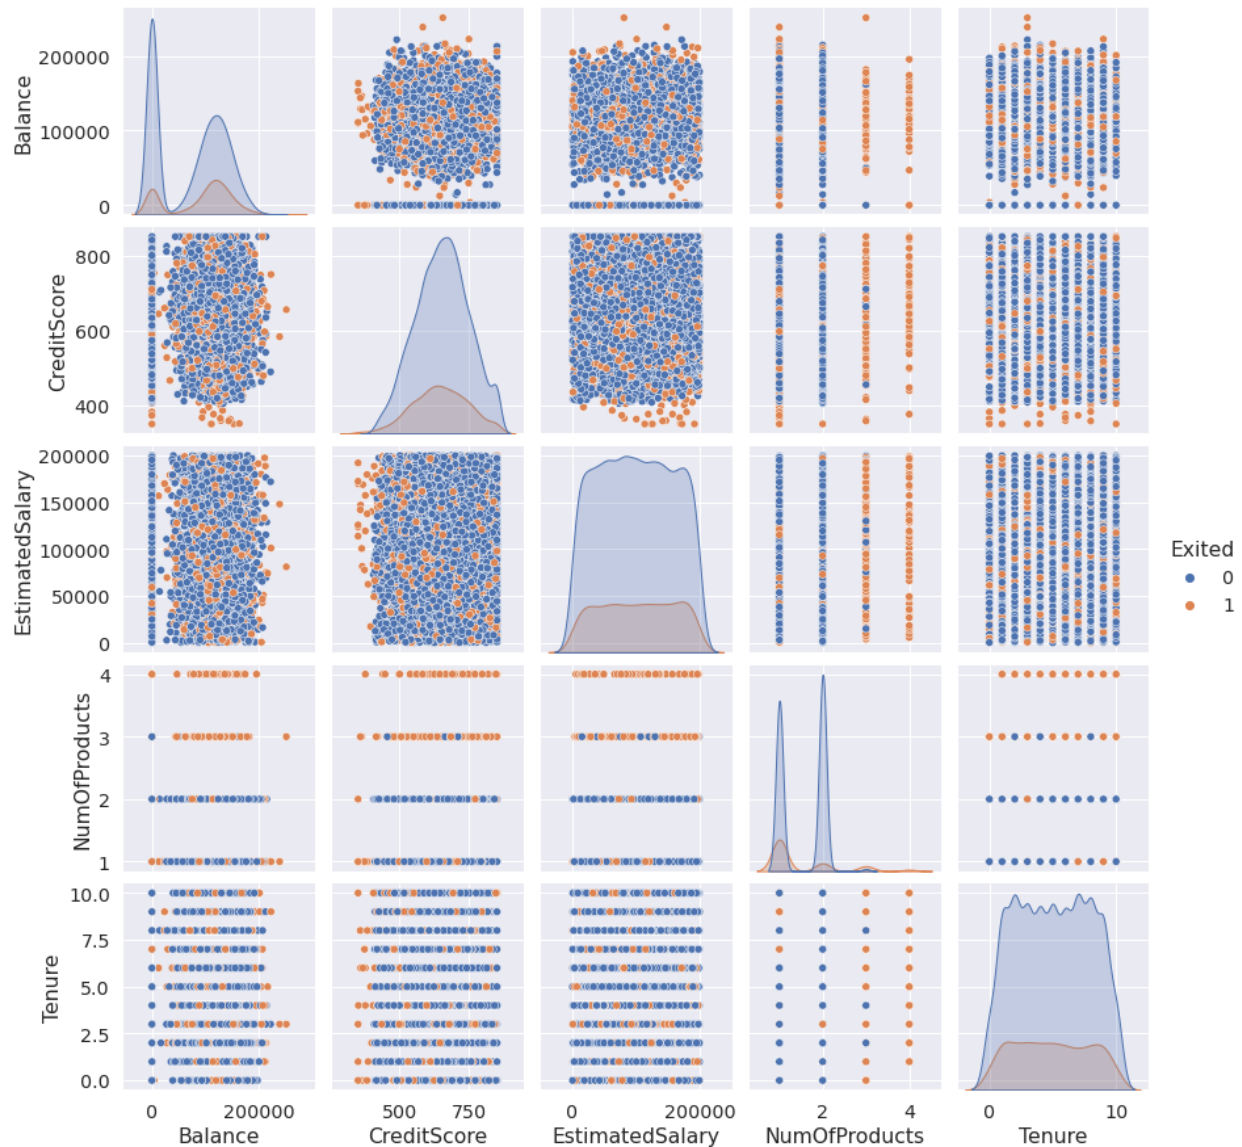
```
plt.figure(figsize=(4,4))
```

```
sns.pairplot(data=df[["Balance","CreditScore","EstimatedSalary","NumOfProducts","Tenure","Exited"]],  
hue="Exited")
```

Out[35]:

<seaborn.axisgrid.PairGrid at 0x7fa0b00a1b10>

<Figure size 288x288 with 0 Axes>



#Perform Descriptive Statistics

```
df=pd.DataFrame(df)
```

```
print(df.sum())
```


CreditScore	6505288
Geography	FranceSpainFranceFranceSpainSpainFranceGermany...
Gender	FemaleFemaleFemaleFemaleFemaleMaleMaleFemaleMa...
Age	389218
Tenure	50128
Balance	764858892.88
NumOfProducts	15302
HasCrCard	7055
IsActiveMember	5151
EstimatedSalary	1000902398.81
Exited	2037

dtype: object

In [39]:

#Perform Descriptive Statistics

```
print("----Sum Value----")
```

```
print(df.sum(1))
```

```
print("-----")
```

```
print("----Product Value----")
```

```
print(df.prod())
```

```
print("-----")
```

----Sum Value----

0 102015.88

1 197002.44

2 274149.37

3 94567.63

4 205492.92

...

9995 97088.64

9996 159633.38

9997 42840.58

9998 168784.83

9999 169159.57

Length: 10000, dtype: float64

-----Product Value-----

CreditScore 0.0

Age 0.0

Tenure 0.0

Balance 0.0

NumOfProducts 0.0

HasCrCard 0.0

IsActiveMember 0.0

EstimatedSalary inf

Exited 0.0

dtype: float64

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:3: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

This is separate from the ipykernel package so we can avoid doing imports until

/usr/local/lib/python3.7/dist-packages/numpy/core/_methods.py:52: RuntimeWarning: overflow encountered in reduce

return umr_prod(a, axis, dtype, out, keepdims, initial, where)

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:6: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

In [38]:

#Perform Descriptive Statistics

```
print("-----Mean Value-----")
```

```
print(df.mean())
```

```
print("-----")
```

```
print("-----Median Value-----")
```

```
print(df.median())
```

```
print("-----")
```

```
print("-----Mode Value-----")
```

```
print(df.mode())
```

```
print("-----")
```

```
-----Mean Value-----
```

CreditScore	650.528800
-------------	------------

Age	38.921800
-----	-----------

Tenure	5.012800
--------	----------

Balance	76485.889288
---------	--------------

NumOfProducts	1.530200
---------------	----------

HasCrCard	0.705500
-----------	----------

IsActiveMember	0.515100
----------------	----------

EstimatedSalary	100090.239881
-----------------	---------------

Exited	0.203700
--------	----------

```
dtype: float64
```

```
-----
```

```
-----Median Value-----
```

CreditScore	652.000
-------------	---------

Age	37.000
-----	--------

Tenure	5.000
--------	-------

Balance	97198.540
---------	-----------

NumOfProducts	1.000
---------------	-------

HasCrCard	1.000
-----------	-------

IsActiveMember	1.000
----------------	-------

[illegible]

[illegible]

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMem
...
9995	True	True	True	True	True	True	True	True	True
9996	True	True	True	True	True	True	True	True	True
9997	True	True	True	True	True	True	True	True	True
9998	True	True	True	True	True	True	True	True	True
9999	True	True	True	True	True	True	True	True	True

10000 rows × 11 columns

In [43]:

```
#Find outliers & replace the outliers
```

```
sns.boxplot(df['Balance'])
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[43]:

<matplotlib.axes._subplots.AxesSubplot at 0x7fa0af6dcf90>

```
#Find outliers & replace the outliers
```

```
print(np.where(df['Balance']>100000))
```

(array([2, 4, 5, ..., 9987, 9993, 9999]),)

9995	771	France	0	39	5	0.00	2	1	0	96270.64	0
9996	516	France	0	35	10	57369.61	1	1	1	101699.77	0
9997	709	France	0	36	7	0.00	1	0	1	42085.58	1

9998	772	Germany	1	42	3	75075.31	2	1	0	92888.52	1
------	-----	---------	---	----	---	----------	---	---	---	----------	---

9999	792	France	0	28	4	130142.79	1	1	0	38190.78	0
------	-----	--------	---	----	---	-----------	---	---	---	----------	---

10000 rows x 11 columns

3	39	1	0.00
---	----	---	------

4	43	2	125510.82
---	----	---	-----------

... ..

9995	39	5	0.00
------	----	---	------

9996	35	10	57369.61
------	----	----	----------

9997	36	7	0.00
------	----	---	------

9998	42	3	75075.31
------	----	---	----------

9999	28	4	130142.79
------	----	---	-----------

[10000 rows x 3 columns]

-----Independent Variables-----

0	1
---	---

1	1
---	---

2	3
---	---

3	2
---	---

4	1
---	---

..

9995	2
------	---

9996	1
------	---

9997	1
------	---

9998	2
------	---

9999	1
------	---

Name: NumOfProducts, Length: 10000, dtype: int64

In []:

#Scale the independent Variables

from sklearn.preprocessing **import** StandardScaler

object= StandardScaler()

standardization

scale = object.fit_transform(df)

print(scale)

```
[[ -0.32622142  0.29351742 -1.04175968 ...  0.97024255  0.02188649
   1.97716468]
 [ -0.44003595  0.19816383 -1.38753759 ...  0.97024255  0.21653375
  -0.50577476]
 [ -1.53679418  0.29351742  1.03290776 ... -1.03067011  0.2406869
   1.97716468]
 ...
 [  0.60498839 -0.27860412  0.68712986 ...  0.97024255 -1.00864308
   1.97716468]
 [  1.25683526  0.29351742 -0.69598177 ... -1.03067011 -0.12523071
   1.97716468]
 [  1.46377078 -1.04143285 -0.35020386 ... -1.03067011 -1.07636976
  -0.50577476]]
```

In []:

#Split the data into training & testing

from sklearn.model_selection **import** train_test_split

In []:

#Split the data into training & testing

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=4, random_state=4)

x_train

Out[]:

	const	EstimatedSalary
--	-------	-----------------

2558	1.0	137903.54
------	-----	-----------

7642	1.0	121765.00
------	-----	-----------

8912	1.0	109470.34
------	-----	-----------

3319	1.0	2923.61
------	-----	---------

6852	1.0	7312.25
------	-----	---------

...
-----	-----	-----

456	1.0	7666.73
-----	-----	---------

6017	1.0	9085.00
------	-----	---------

709	1.0	147794.63
-----	-----	-----------

8366	1.0	102515.42
------	-----	-----------

1146	1.0	54776.64
------	-----	----------

9996 rows × 2 columns

In []:

#Split the data into training & testing

x_test

Out[]:

	const	EstimatedSalary
--	-------	-----------------

1603	1.0	23305.85
------	-----	----------

8713	1.0	41248.80
------	-----	----------

	const	EstimatedSalary
--	-------	-----------------

4561	1.0	143317.42
------	-----	-----------

6600	1.0	174123.16
------	-----	-----------

In []:

#Split the data into training & testing

y_train

Out[]:

2558 727

7642 811

8912 623

3319 430

6852 600

...

456 733

6017 487

709 686

8366 637

1146 614

Name: CreditScore, Length: 9996, dtype: int64

In []:

#Split the data into training & testing

y_test

Out[]:

1603 576

8713 786

4561 562

6600 505

Name: CreditScore, dtype: int64

